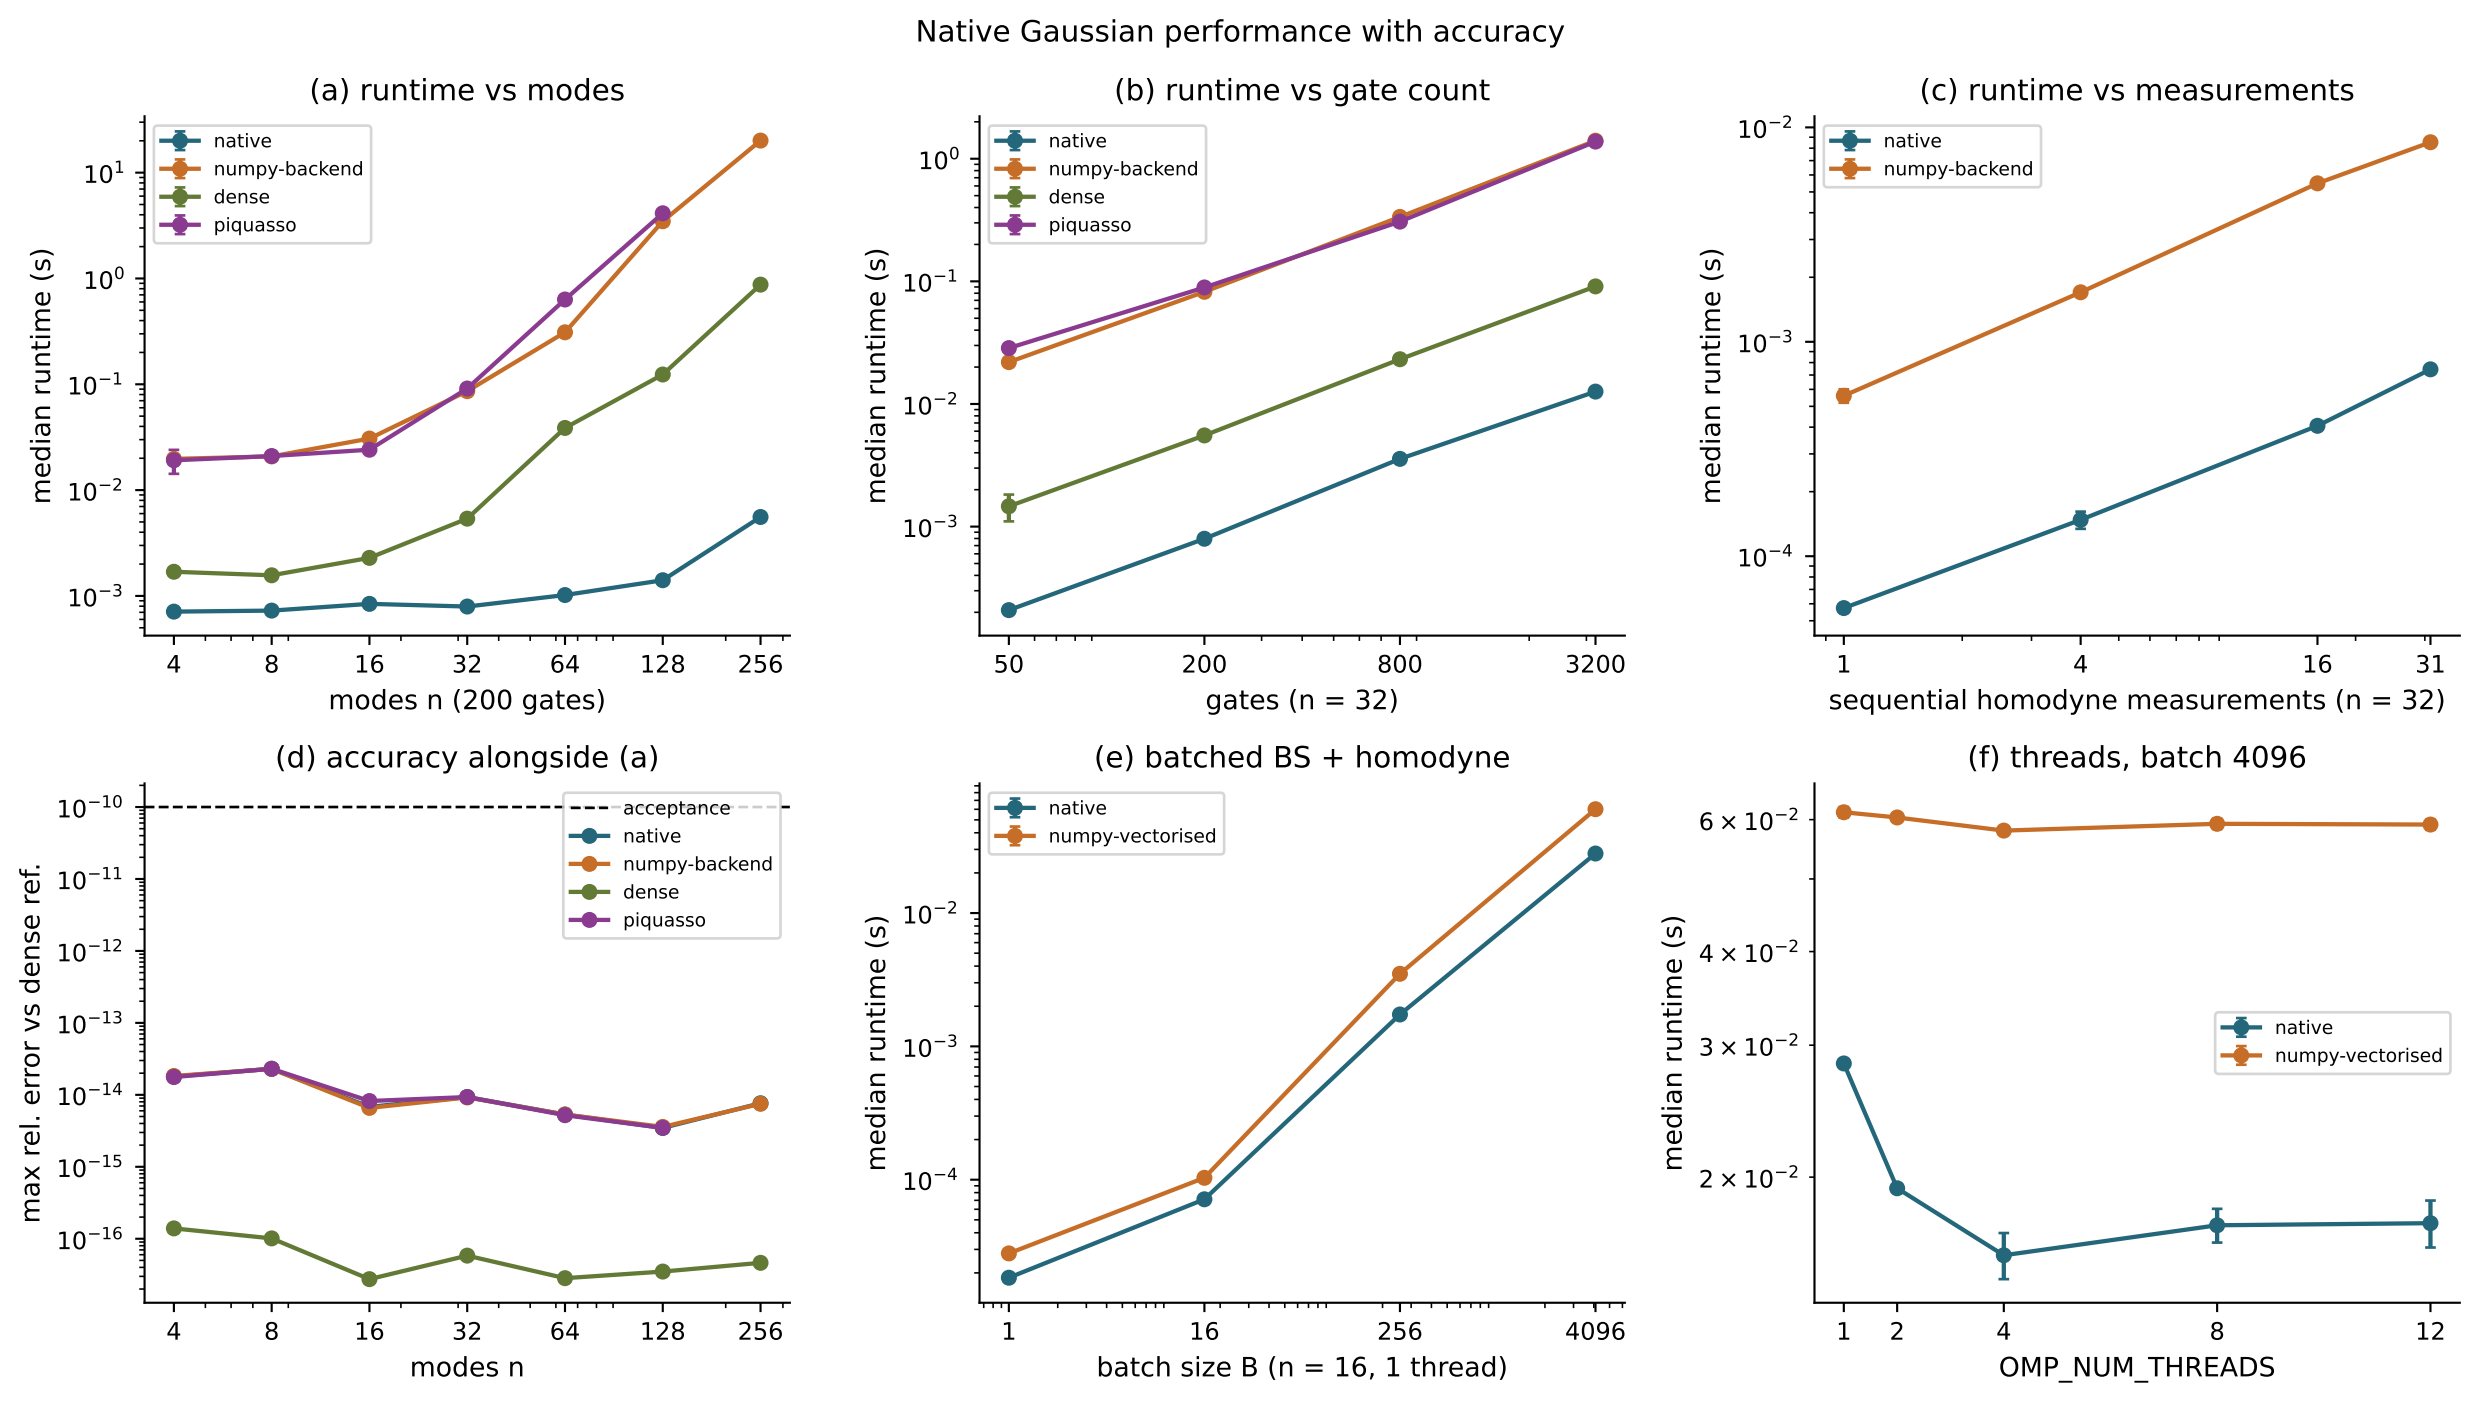
Native Gaussian performance with accuracy
(a) runtime vs modes
native
numpy-backend
dense
piquasso
median runtime (s)
modes n (200 gates)
(b) runtime vs gate count
gates (n = 32)
(c) runtime vs measurements
sequential homodyne measurements (n = 32)
(d) accuracy alongside (a)
acceptance
max rel. error vs dense ref.
modes n
(e) batched BS + homodyne
numpy-vectorised
batch size B (n = 16, 1 thread)
(f) threads, batch 4096
OMP_NUM_THREADS